



```

    "std": 1, "min": 0, "max": 8, "n
    "num_unique_values": 7, "samples": [1,
0], "semantic_type": "\", "column": "Parch", "properties": {"dtype": "number", "std": 0, "min": 0, "max": 6, "n
    "num_unique_values": 7, "samples": [0,
1], "semantic_type": "\", "column": "Ticket", "properties": {"dtype": "string", "n
    "num_unique_values": 681, "samples": ["11774", "248740"], "semantic_type": "\", "description": "\n      \", "column": "Fare", "properties": {"n
    "dtype": "number", "std": 49.693428597180905, "min": 0.0, "max": 512.3292, "n
    "num_unique_values": 248, "samples": [11.2417, 51.8625], "semantic_type": "\", "description": "\n      \", "column": "Cabin", "properties": {"n
    "dtype": "category", "n
    "num_unique_values": 147, "samples": ["D45", "B49"], "semantic_type": "\", "description": "\n      \", "column": "Embarked", "properties": {"n
    "dtype": "category", "n
    "num_unique_values": 3, "samples": ["S", "C"], "semantic_type": "\", "description": "\n      \", "n
    }], "type": "dataframe", "variable_name": "df"}}

# @title Understand the Dataset
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
 #   Column      Non-Null Count Dtype  
--- 
 0   PassengerId 891 non-null   int64  
 1   Survived     891 non-null   int64  
 2   Pclass       891 non-null   int64  
 3   Name         891 non-null   object  
 4   Sex          891 non-null   object  
 5   Age          714 non-null   float64 
 6   SibSp        891 non-null   int64  
 7   Parch        891 non-null   int64  
 8   Ticket       891 non-null   object  
 9   Fare          891 non-null   float64 
 10  Cabin         204 non-null   object  
 11  Embarked     889 non-null   object 

```

```
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB

df.describe()

{"summary": {"name": "df", "rows": 8, "fields": [{"column": "PassengerId", "properties": {"dtype": "number", "std": 320.8159711429856, "min": 1.0, "max": 891.0, "num_unique_values": 6, "samples": [891.0, 446.0, 668.5], "semantic_type": "\\", "description": "\n"}, {"column": "Survived", "properties": {"dtype": "number", "std": 314.8713661874558, "min": 0.0, "max": 891.0, "num_unique_values": 5, "samples": [0.3838383838383838, 1.0, 0.4865924542648585], "semantic_type": "\\", "description": "\n"}, {"column": "Pclass", "properties": {"dtype": "number", "std": 314.2523437079693, "min": 0.8360712409770513, "max": 891.0, "num_unique_values": 6, "samples": [2.308641975308642, 3.0], "semantic_type": "\\", "description": "\n"}, {"column": "Age", "properties": {"dtype": "number", "std": 242.9056731818781, "min": 0.42, "max": 714.0, "num_unique_values": 8, "samples": [29.69911764705882, 28.0, 714.0], "semantic_type": "\\", "description": "\n"}, {"column": "SibSp", "properties": {"dtype": "number", "std": 314.4908277465442, "min": 0.0, "max": 891.0, "num_unique_values": 6, "samples": [8.0, 0.5230078563411896, 0.38159371492704824], "semantic_type": "\\", "description": "\n"}, {"column": "Parch", "properties": {"dtype": "number", "std": 314.65971717879, "min": 0.0, "max": 891.0, "num_unique_values": 5, "samples": [0.8060572211299559], "semantic_type": "\\", "description": "\n"}, {"column": "Fare", "properties": {"dtype": "number", "std": 330.6256632228577, "min": 0.0, "max": 891.0, "num_unique_values": 8, "samples": [32.204207968574636, 14.4542, 891.0], "semantic_type": "\\", "description": "\n"}]}, "type": "dataframe"}}

df.isnull().sum()
```

```
PassengerId      0
Survived         0
Pclass           0
Name             0
Sex              0
Age             177
SibSp            0
Parch            0
Ticket          0
Fare             0
Cabin          687
Embarked        2
dtype: int64

# @title Data Cleaning
df['Age'].fillna(df['Age'].mean(), inplace=True)

/tmp/ipython-input-1226013979.py:1: FutureWarning: A value is trying
to be set on a copy of a DataFrame or Series through chained
assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never
work because the intermediate object on which we are setting values
always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try
using 'df.method({col: value}, inplace=True)' or df[col] =
df[col].method(value) instead, to perform the operation inplace on the
original object.

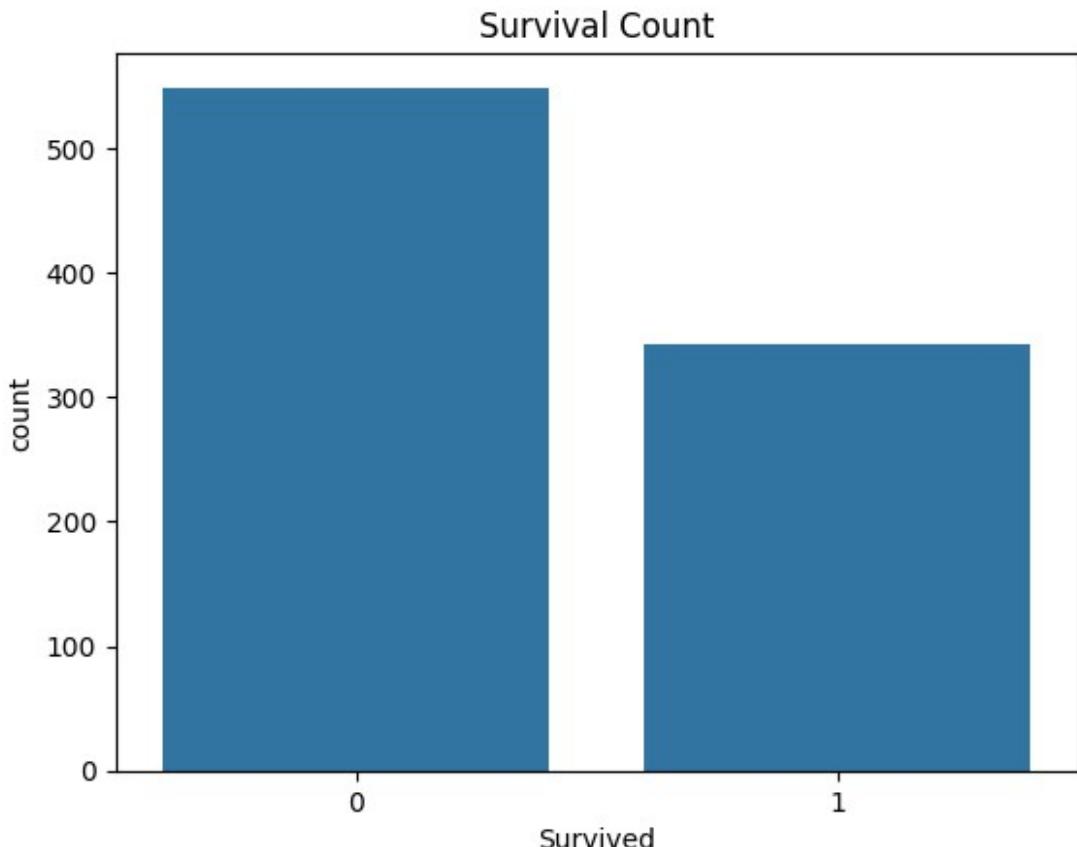
df['Age'].fillna(df['Age'].mean(), inplace=True)
df['Embarked'].fillna(df['Embarked'].mode()[0], inplace=True)

/tmp/ipython-input-411043493.py:1: FutureWarning: A value is trying to
be set on a copy of a DataFrame or Series through chained assignment
using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never
work because the intermediate object on which we are setting values
always behaves as a copy.

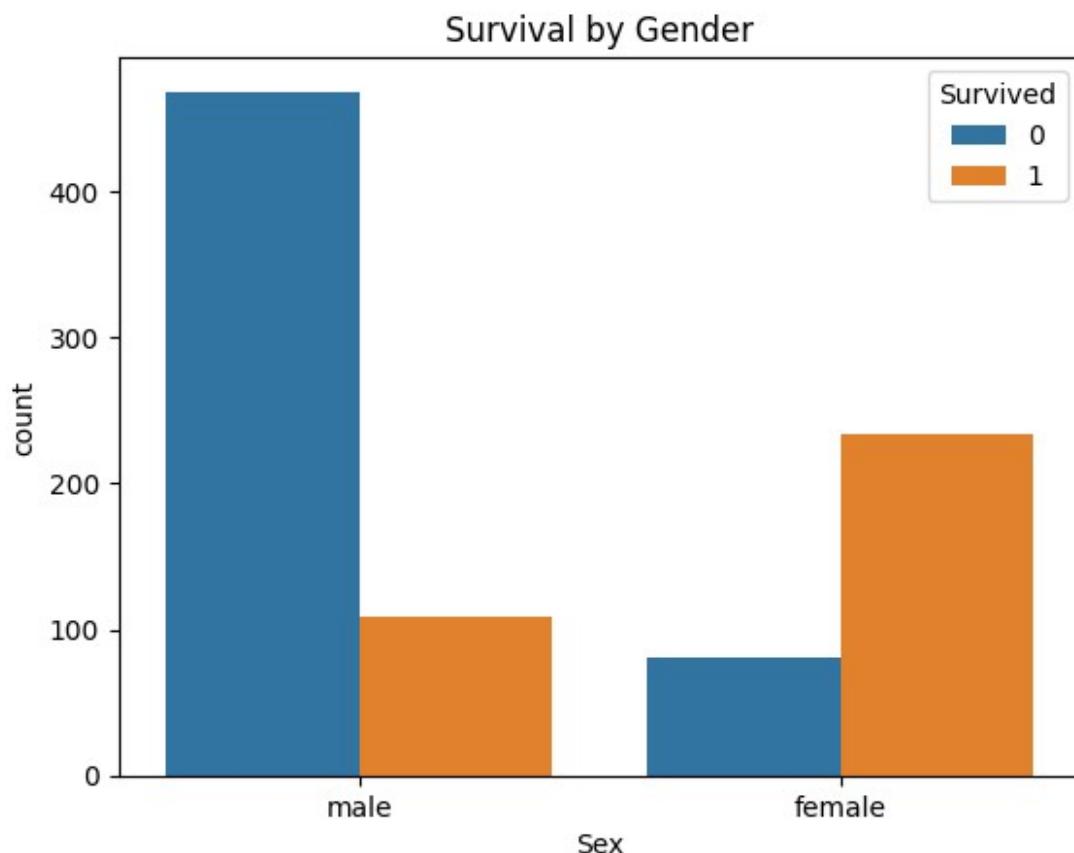
For example, when doing 'df[col].method(value, inplace=True)', try
using 'df.method({col: value}, inplace=True)' or df[col] =
df[col].method(value) instead, to perform the operation inplace on the
original object.

df['Embarked'].fillna(df['Embarked'].mode()[0], inplace=True)
df.drop(columns=['Cabin'], inplace=True)
```

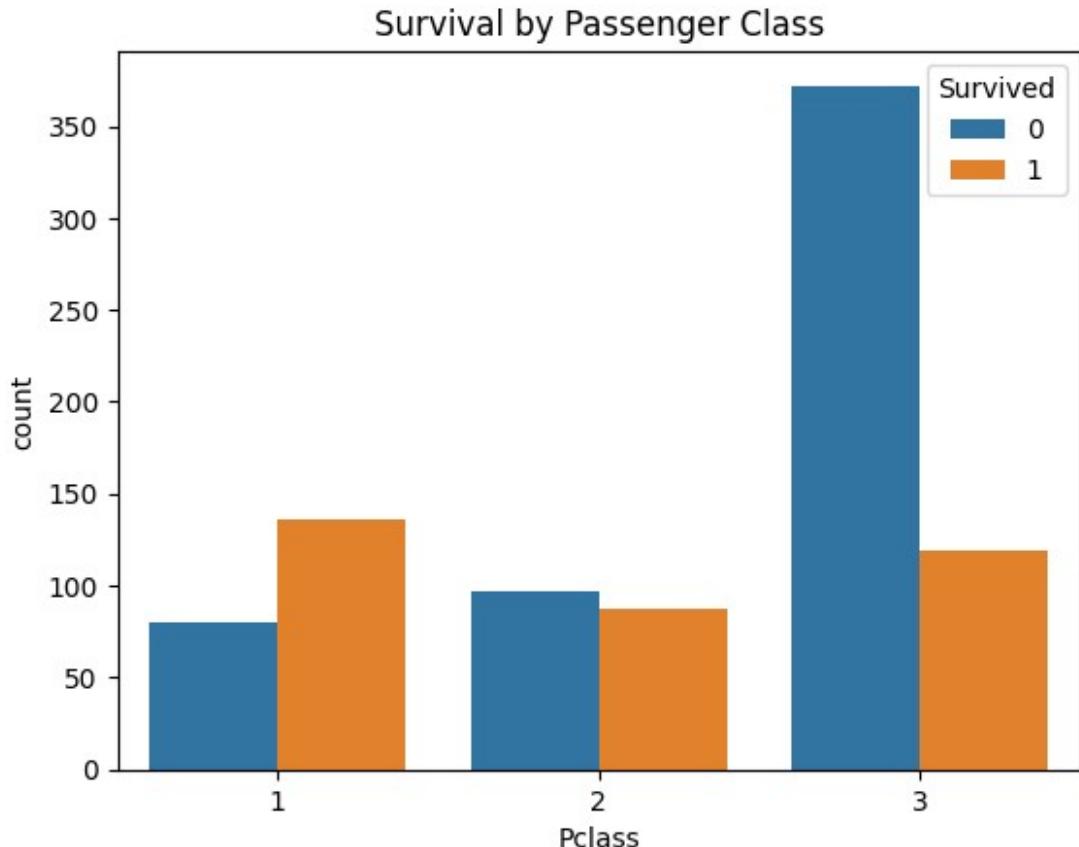
```
# @title Exploratory Data Analysis (EDA)
sns.countplot(x='Survived', data=df)
plt.title("Survival Count")
plt.show()
```



```
sns.countplot(x='Sex', hue='Survived', data=df)
plt.title("Survival by Gender")
plt.show()
```



```
sns.countplot(x='Pclass', hue='Survived', data=df)
plt.title("Survival by Passenger Class")
plt.show()
```



```

# @title Encode Categorical Variables
le = LabelEncoder()

df['Sex'] = le.fit_transform(df['Sex'])
df['Embarked'] = le.fit_transform(df['Embarked'])

# @title Feature Selection
X = df[['Pclass', 'Sex', 'Age', 'SibSp', 'Parch', 'Fare', 'Embarked']]
y = df['Survived']

# @title Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

# @title Build the Model
model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)

LogisticRegression(max_iter=1000)

# @title Model Prediction
y_pred = model.predict(X_test)

```

```

# @title Model Evaluation
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)

Accuracy: 0.8100558659217877

print(classification_report(y_test, y_pred))

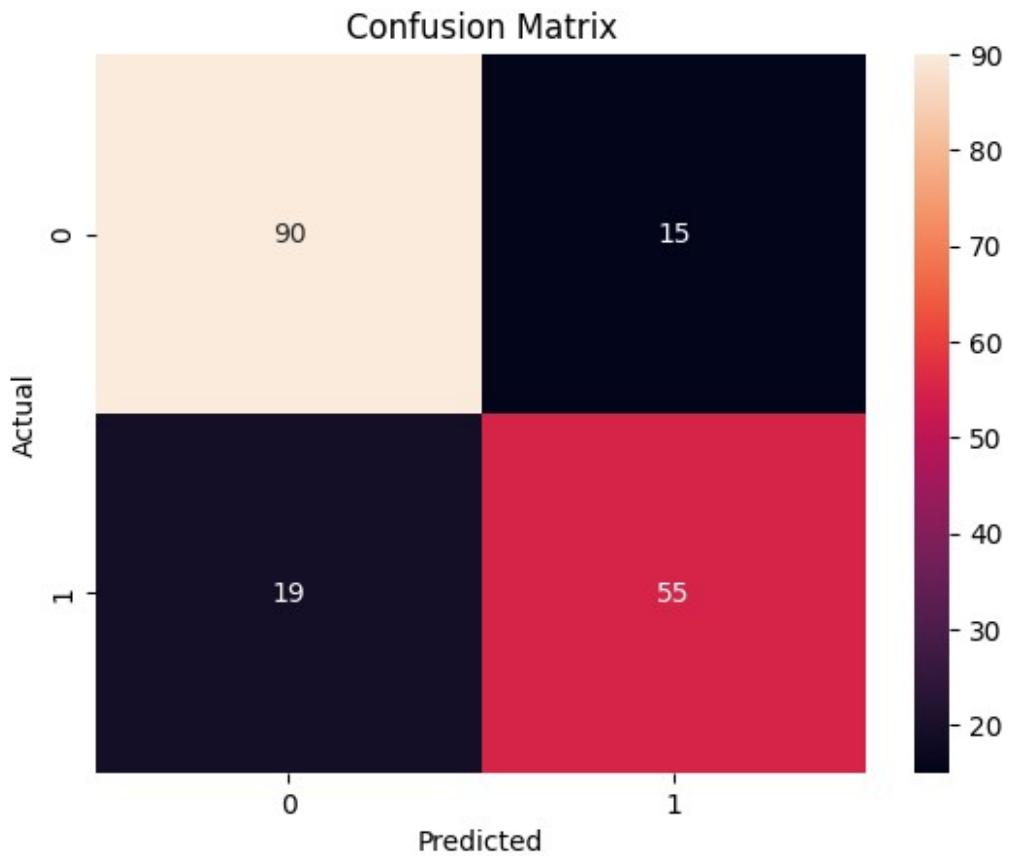
      precision    recall  f1-score   support

          0       0.83      0.86      0.84     105
          1       0.79      0.74      0.76      74

   accuracy                           0.81      179
  macro avg       0.81      0.80      0.80      179
weighted avg       0.81      0.81      0.81      179

cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d')
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

```



```
# @title Conclusion
print("Logistic Regression model successfully predicts Titanic
survival.")
print("Gender, Passenger Class, and Fare play major roles in
survival.")
```

```
Logistic Regression model successfully predicts Titanic survival.
Gender, Passenger Class, and Fare play major roles in survival.
```